

yoavg / cnn

forked from clab/cnn

Watch

2

Star

13

Fork

53

<> Code

Pull requests

0

Pulse

Graphs

C++ neural network library

517 commits

3 branches

0 releases

14 contributors

Branch: master

New pull request

New file

Find file

HTTPS

https://github.com/yoavg/cnn.

Download ZIP

This branch is 39 commits behind clab:master.

Pull request

Compare

yoavg pass cmdline args to pycnn init

Latest commit 6e49f0c 27 days ago

cmake	test stub	4 months ago
cnn	Merge branch 'master' of github.com:clab/cnn	28 days ago
examples	bad cmake dependencies	a month ago
external	skip-gram RNN language model	7 months ago
pycnn	pass cmdline args to pycnn init	27 days ago
pyexamples	fixed embarassing mistake in RNNs.ipynb	4 months ago
tests	fix test failure	a month ago
.gitignore	rename HSM CFMS	2 months ago
.gitmodules	skip-gram RNN language model	7 months ago
.travis.yml	CI gcc 4.8	3 months ago
CMakeLists.txt	Merge branch 'master' of github.com:clab/cnn	a month ago
INSTALL.md	INSTALL.md	5 months ago
LICENSE	lic terms	7 months ago
README.md	README fix to use Expression notation and fix bugs	3 months ago
TODO.cnn	Update TODO.cnn	3 months ago
config.h.cmake	multibackend support, eigen cleanup and optimizations	9 months ago

README.md

cnn

C++ neural network library

Getting started

You need the [development version of the Eigen library](#) for this software to function properly. If you use the current stable release, you will get an error like the following:

```
Assertion failed: (false && "heap allocation is forbidden (EIGEN_NO_MALLOC is defined)", function check
```

Building

First you need to fetch the dependent libraries

```
git submodule init
git submodule update
```

In `src`, you need to first use `cmake` to generate the makefiles

```
mkdir build
cd build
cmake .. -DEIGEN3_INCLUDE_DIR=/path/to/eigen
```

Then to compile, run

```
make -j 2
```

To see that things have built properly, you can run

```
./examples/xor
```

which will train a multilayer perceptron to predict the xor function.

Building without Eigen installed

If you don't have Eigen installed, the instructions below will fetch and compile both `Eigen` and `cnn`.

```
git clone https://github.com/clab/cnn.git
hg clone https://bitbucket.org/eigen/eigen/

cd cnn/
mkdir build
cd build
cmake .. -DEIGEN3_INCLUDE_DIR=../eigen
make -j 2
```

Debugging

If you want to see the compile commands that are used, you can run

```
make VERBOSE=1
```

Training Models

An illustration of how models are trained (for a simple logistic regression model) is below:

```
// *** First, we set up the structure of the model
// Create a model, and an SGD trainer to update its parameters.
Model mod;
SimpleSGDTrainer sgd(&mod);
// Create a "computation graph," which will define the flow of information.
ComputationGraph cg;
// Initialize a 1x3 parameter vector, and add the parameters to be part of the
// computation graph.
Expression W = parameter(cg, mod.add_parameters({1, 3}));
// Create variables defining the input and output of the regression, and load them
// into the computation graph. Note that we don't need to set concrete values yet.
vector<cnn::real> x_values(3);
Expression x = input(cg, {3}, &x_values);
cnn::real y_value;
Expression y = input(cg, &y_value);
// Next, set up the structure to multiply the input by the weight vector, then run
// the output of this through a logistic sigmoid function (logistic regression).
Expression y_pred = logistic(W*x);
// Finally, we create a function to calculate the loss. The model will be optimized
```

```
// to minimize the value of the final function in the computation graph.
Expression l = binary_log_loss(y_pred, y);
// We are now done setting up the graph, and we can print out its structure:
cg.PrintGraphviz();

// *** Now, we perform a parameter update for a single example.
// Set the input/output to the values specified by the training data:
x_values = {0.5, 0.3, 0.7};
y_value = 1.0;
// "forward" propagates values forward through the computation graph, and returns
// the loss.
cnn::real loss = as_scalar(cg.forward());
// "backward" performs back-propagation, and accumulates the gradients of the
// parameters within the "Model" data structure.
cg.backward();
// "sgd.update" updates parameters of the model that was passed to its constructor.
// Here 1.0 is the scaling factor that allows us to control the size of the update.
sgd.update(1.0);
```

Note that this very simple example that doesn't cover things like memory initialization, reading/writing models, recurrent/LSTM networks, or adding biases to functions. The best way to get an idea of how to use cnn for real is to look in the `example` directory, particularly starting with the simplest `xor` example.

